

## Lecture 9: Bayesian Hierarchical Models 2

Lecturer: Ramesh Sridharan

In the last lecture, we introduced the idea of Gaussian mixture models (GMMs). In this lecture, we will introduce the *Expectation-Maximization* Algorithm as a way of performing unsupervised learning to learn GMMs from data.

## 1 Gaussian Mixture Models

Suppose we have a random variable  $Y$  with an unknown distribution  $\mathbb{P}(Y)$  that may not fit into any known class of distributions that we are familiar with. However, we would like to come up with a model of the distribution that we can interpret and sample from easily, that closely matches the distribution of the random variable  $Y$ . One common model that is simple to use yet powerful enough to represent complex distributions are mixtures of Gaussians as illustrated in Figure 9.1.




Figure 9.1: The probability density function of  $Y$  and the GMM that describes it.

To model  $\mathbb{P}(Y)$  as a mixture of Gaussians we assume that there is more structure underlying the random variable. In particular, we assume that there is a hidden random variable  $X$  taking values in  $i = 1, 2, \dots, d$  such that:

$$\mathbb{P}(Y|X = i) = \mathcal{N}(\mu_i, \sigma_i^2)$$

Thus,  $\mathbb{P}(Y) = \sum_{i=1}^d \pi_i \mathcal{N}(Y; \mu_i, \sigma_i^2)$ , where  $\pi_i = \mathbb{P}(X = i)$ , and we model  $\mathbb{P}(Y)$  as a mixture of  $d$  Gaussians.

**Example 9.1.** For the distribution showed in Figure 9.1,  $\pi_1 = 0.1, \pi_2 = 0.5, \pi_3 = 0.2, \pi_4 = 0.2$ .

Looking at GMMs from a Bayesian perspective, one can think of  $\pi_i$  as the prevalence that  $Y$  came from the normal with mean  $\mu_i$  and variance  $\sigma_i^2$ .

## 2 Expectation-Maximization

In the previous section we discussed Gaussian Mixture models and their ability to model complex distributions. In this section we will present a method of learning the parameters of a GMM  $\mu_i, \sigma_i^2$ , and  $\pi_i$  from data. One of the biggest problems with learning Gaussian mixture models is that simply maximizing the likelihood of the data does not work since we typically only observe  $y_1, \dots, y_n$ , and not the hidden variables  $x_1, \dots, x_n$ . Indeed letting,  $\theta_i = (\mu_i, \sigma_i^2)$  be the parameters of Gaussian  $i$ , we can write the likelihood of the data points  $y_j$  as:

$$\mathbb{P}(y_j | \theta_1, \dots, \theta_n) = \prod_{i=1}^d (\pi_i \mathcal{N}(y_j; \mu_i, \sigma_i^2))^{\mathbb{I}(x_j=i)}$$

The log likelihood of all of the data is therefore:

$$\ell(y; \theta_1, \dots, \theta_n) = \sum_{j=1}^n \sum_{i=1}^d \mathbb{I}(x_j = i) (\log(\pi_i) + \log(\mathcal{N}(y_j; \mu_i, \sigma_i^2)))$$

To maximize the log-likelihood of the data, we therefore need to find:

1. The values of the hidden variable  $x_j$  for each data point  $y_j$  for  $j = 1, \dots, n$ .
2. The mean  $\mu$ , variance  $\sigma^2$ , and prevalence  $\pi$  associated with each of the  $d$ .

This is not straightforward. If we knew which of the data points  $y_j$  came from which Normal distribution  $\mathcal{N}(\mu_i, \sigma_i^2)$  (i.e. we know  $x_i$  for each  $y_i$ ), we could find the maximum likelihood estimates of  $\mu_i$ ,  $\sigma_i^2$ , and  $\pi_i$ . Further, if we knew the values of  $\pi_i$  as well as  $\mu_i$  and  $\sigma_i^2$  we could find the posterior distribution over  $X_j$  for each data point  $y_j$ . This posterior is given by:

$$\mathbb{P}(X_j = i | y_j) = \frac{\pi_i \mathcal{N}(y_j; \mu_i, \sigma_i^2)}{\sum_{k=1}^d \pi_k \mathcal{N}(y_j; \mu_k, \sigma_k^2)}$$

Therefore, we have a chicken and egg scenario. Given data points  $y_1, \dots, y_n$ , if we knew  $x_1, \dots, x_n$  we could find the values of the parameters of our GMM,  $(\pi_i, \mu_i, \sigma_i^2)$  for  $i = 1, \dots, d$ , and if we knew the values of the parameters of our GMM,  $(\pi_i, \mu_i, \sigma_i^2)$  for  $i = 1, \dots, d$ , we could predict  $x_1, \dots, x_n$ .

To solve this problem, we introduce the *Expectation-Maximization Algorithm* or EM algorithm for short. There are 3 main ideas behind the EM Algorithm:

1. Randomly initialize  $\theta_i$  and  $\pi_i$ .

2. Given fixed  $\theta_i$  and  $\pi_i$ , for each data point  $y_j$  approximate the probability that  $y_j$  comes from Gaussian  $i$ , denoted  $Z_j(i) = \mathbb{P}(X_j = i|y_j)$ .
3. Given fixed distributions  $Z_j$  find the values of  $\theta_i$  and  $\pi_i$  that maximize the expected likelihood of the data (over the distributions  $Z_j(i)$ ):

$$\theta^* = \arg \max_{\theta_1, \dots, \theta_d} \mathbb{E}_Z[\ell(y; \theta_1, \dots, \theta_d)]$$

Since  $\mathbb{E}[\mathbb{I}(x_j = i)] = \Pr(X_j = i|Y_i) = Z_j(i)$ , this simplifies to:

$$\theta^* = \arg \max_{\theta_1, \dots, \theta_d} \sum_{j=1}^n \sum_{i=1}^d Z_j(i) (\log(\pi_i) + \log(\mathcal{N}(y_j; \mu_i, \sigma_i^2)))$$

4. Iterate between the two sub-problems until convergence.

**Remark 9.2.** The EM algorithm can be shown to maximize the lower bound on the log-likelihood of the data at each iteration, meaning that as the algorithm runs we have more and more confidence that the log-likelihood of the data is improving.

We now outline the *EM* algorithm with unknown  $\mu_i$ ,  $\sigma_i^2$ ,  $\pi_i$  for a mixture of  $d$  Gaussians given  $y_1, \dots, y_n$ .

---

**Algorithm 1** Expectation-Maximization Algorithm for Gaussian Mixture Models

---

**Input:** Data:  $y_1, \dots, y_n$ , Number of Gaussians in the mixture  $d$ , number of iterations  $r$

**Output:**  $(\pi_i, \mu_i, \sigma_i^2)$  for  $i = 1, \dots, d$ .

Randomly Initialize  $(\pi_{i,0}, \mu_{i,0}, \sigma_{i,0}^2)$  **for**  $t = 1$  **to**  $r$  **do**

Expectation Step: **for**  $j = 1$  **to**  $n$  **do**

**for**  $i = 1$  **to**  $d$  **do**

$Z_j(i) \leftarrow \frac{\pi_{i,t-1} \mathcal{N}(y_j; \mu_{i,t-1}, \sigma_{i,t-1}^2)}{\sum_{k=1}^d \pi_{k,t-1} \mathcal{N}(y_j; \mu_{k,t-1}, \sigma_{k,t-1}^2)}$

**end**

**end**

Maximization Step: **for**  $i = 1$  **to**  $d$  **do**

$N_{i,t} \leftarrow \sum_{j=1}^n Z_j(i).$

$\mu_{i,t} \leftarrow \frac{1}{N_{i,t}} \sum_{j=1}^n Z_j(i) y_j.$

$\sigma_{i,t} \leftarrow \frac{1}{N_{i,t}} \sum_{j=1}^n Z_j(i) (y_j - \mu_{i,t})^2.$

$\pi_{i,t} \leftarrow \frac{N_{i,t}}{n}.$

**end**

**end**

---

Note that the update for  $\mu_i$  and  $\sigma_i^2$  are both the maximizers of the expected likelihood using the straightforward derivation we have seen in previous lectures and discussions. The update for  $\pi_i$ , however, requires maximizing the expected likelihood while constraining  $\sum_{i=1}^d \pi_i = 1$ . This derivation requires using solving a constrained optimization problem which is outside the scope of this class.

**Remark 9.3.** Note that the EM algorithm can be very sensitive to the initialization, and is not guaranteed to converge to the same solution from any initialization.